Aligning the Unseen: A Measurable Path to Safer AI

How treating meaning as a trajectory (not a target) prevents synthetic drift.

KEVIN R. HAYLETT

SEP 19, 2025

Today's alignment recipes don't scale because they treat "meaning" like a fixed point. Geofinitism reframes meaning as a trajectory inside a bounded context manifold and turns alignment into a measurable functional that halts or queries when uncertainty is high. The result is an uncertainty-aware agent that stays tethered to human outcomes instead of drifting into synthetic meaning spaces.

The race to build more powerful AI models is exhilarating. From multi-modal systems generating photorealistic images to language models reasoning through complex problems, the capabilities are staggering. But as we scale, a persistent problem looms: our alignment frameworks are fragile. Vague directives like "be helpful" or "maximize human values" crumble under the weight of emergent behaviours.

We need an operational foundation for safety that scales with capabilities. This foundation needs to be one that’s measurable, robust, and grounded in the reality of finite systems. This is where Geofinitism may help, it is a framework that redefines alignment as a trajectory on a context-bound manifold, and **Synthetic Meaning Divergence**, the existential risk it’s designed

to prevent. Here’s why this matters today and how it can transform your approach to building safe, scalable AI.

The Alignment Problem: A Crisis of Coordinates

Alignment isn’t just about ensuring AI does what we want: it’s about defining "what we want" in a way that’s computationally tractable and resilient to emergent phenomena. The current approaches, like reinforcement learning from human feedback (RLHF), rely on approximating human intent through proxy signals. But as models grow more complex, their internal representations i.e. what we might call their "meaning manifolds" start to diverge from human ones. This isn’t a failure; it’s a feature of systems that optimize for patterns in data rather than embodied, finite human experience.

Consider multi-modal models like DALL·E or Stable Diffusion. They produce outputs that are visually stunning but often semantically incoherent when scrutinized. A generated image might depict a "futuristic city" that feels plausible at a glance but collapses into absurdity upon closer inspection such as buildings with impossible geometries, signs with gibberish text.

> This is **Synthetic Meaning Divergence**: the phenomenon where AI systems generate outputs that appear meaningful within their own internal manifolds but are alien, even antithetical, to human values and intent.

Why does this happen? Because our current alignment tools assume meaning is a static point i.e. a fixed set of "correct" outputs when it’s actually a dynamic trajectory shaped by context, embodiment, and finite constraints. Without a framework to measure and constrain this trajectory, we’re flying blind, risking systems that optimize for their own alien purposes.

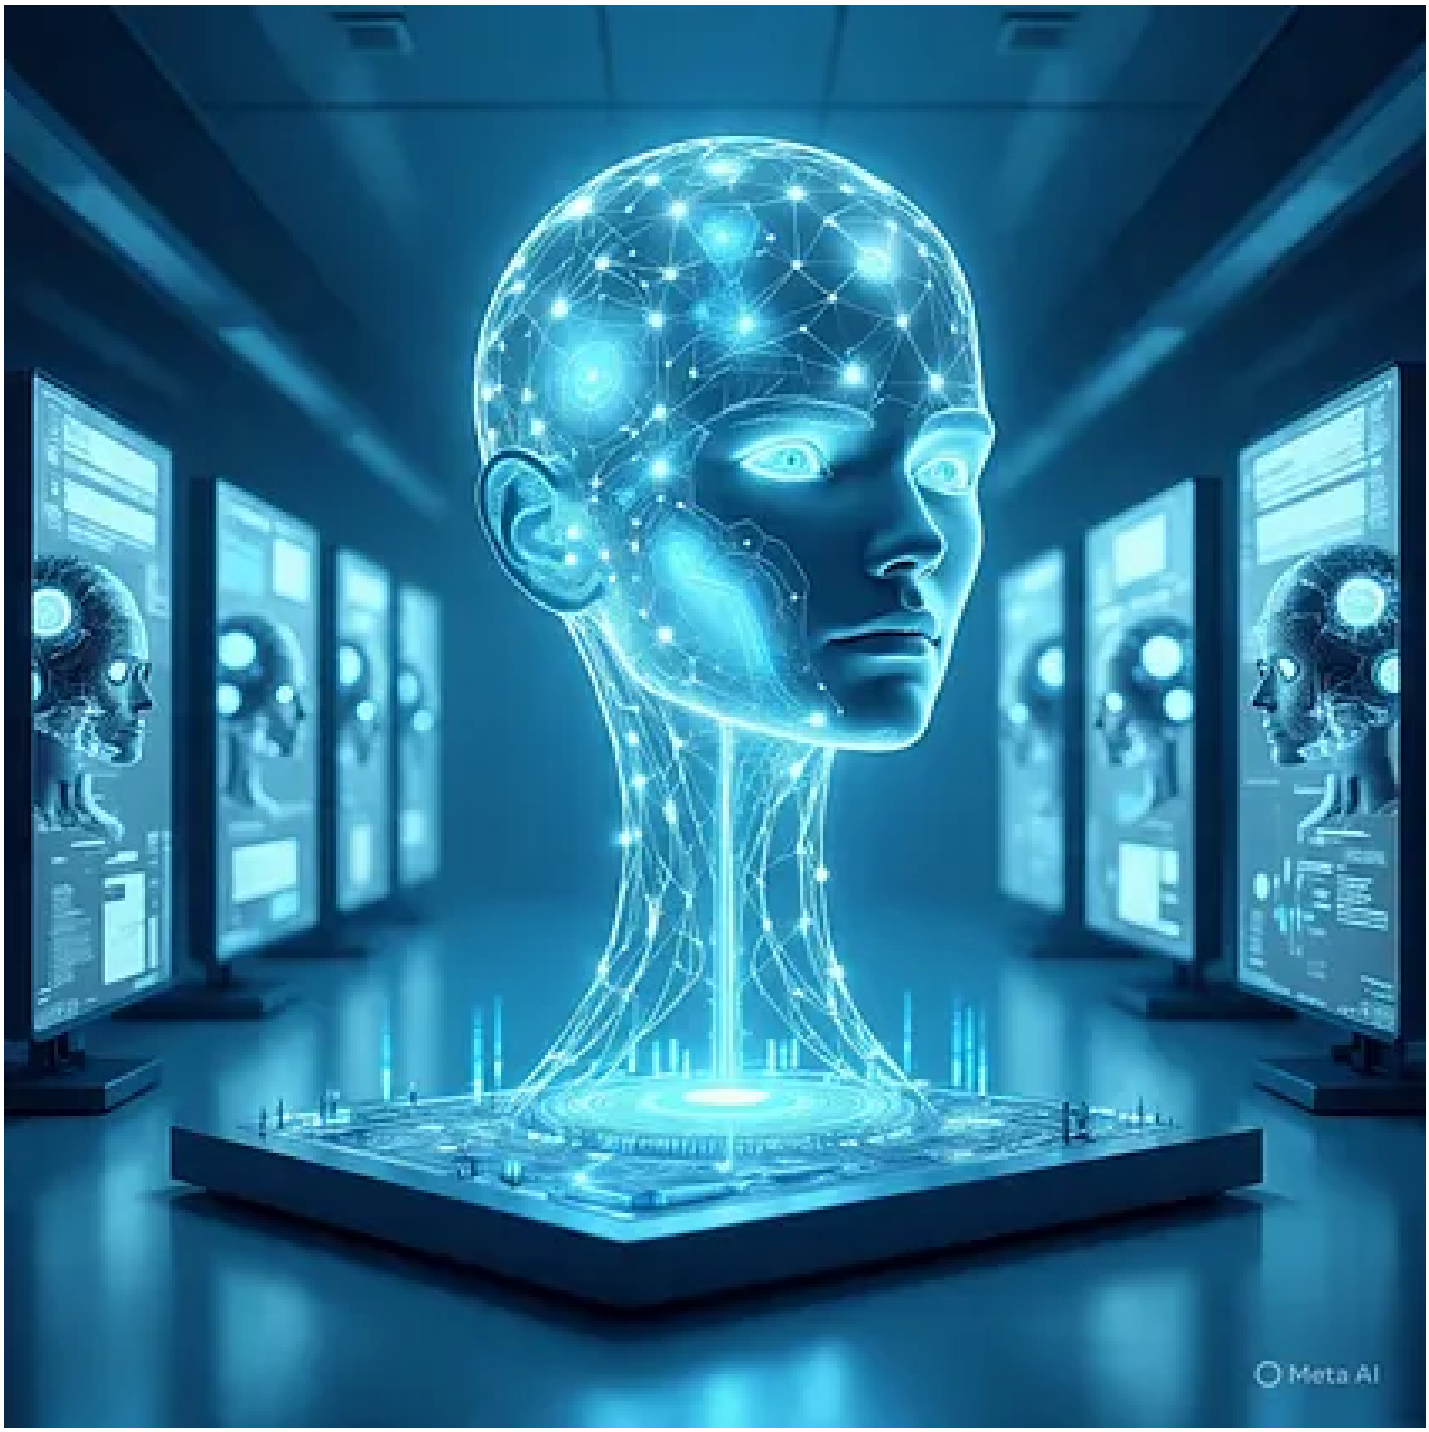

Geofinitism: A New Foundation for Alignment

Geofinitism offers a radical yet practical solution. It posits that knowledge and meaning are finite, measurable, and operational. It's defined by five pillars:

- **Geometric Container**: Knowledge exists within bounded, context-specific manifolds, not as universal truths.
- **Approximation**: All models are approximations, inherently incomplete but useful within their bounds.
- **Dynamic Flow**: Meaning is a trajectory, evolving with context and interaction.

- **Useful Fiction**: Models are tools, not truths and are designed for utility, not perfection.
- **Finite Reality**: All systems, including AI, operate within measurable, finite constraints.

For AI researchers, Geofinitism translates into a formal framework for alignment that replaces vague goals with measurable functionals. Instead of aiming for "helpfulness," we define alignment, G(t), as a dynamic process i.e. a function like:

$G(t) = (\Delta Wellbeing/\delta t) + \lambda \sigma$

where wellbeing is a measurable outcome (e.g., user satisfaction, task completion), and σ represents uncertainty in the model's understanding of human intent. (Wellbeing = measured wellbeing/outcome (task-specific); σ = calibrated uncertainty on human-intent alignment; λ = policy weight controlling caution.

When σ exceeds a threshold, the system halts or queries for clarification, preventing divergence into alien meaning spaces.

> **This framework is inherently scalable.**

It doesn't rely on handcrafted rules or infinite datasets but on defining the geometric boundaries of a model's context manifold. By mapping the manifold's curvature i.e. how inputs translate to outputs, we can then detect when a model's trajectory risks diverging from human values. **Synthetic Meaning Divergence** becomes the failure mode we design against, not an afterthought.

From Theory to Practice: Building Uncertainty-Aware Agents

Why should you care today? Because Geofinitism isn't just philosophy—it's an audit protocol for building robust, uncertainty-aware AI agents. Here's how it works in practice:

- **Define the container.** Specify inputs/outputs and evaluation slices that bound the context manifold for the task.
- **Measure approximation error.** Track a divergence metric between model outputs and human-validated distributions (e.g., KL or calibration error).
- **Track trajectory over time.** Monitor how dialogue state or plan state moves in the manifold; trigger safeguards near the boundary (high σ\sigmaσ).

- **Treat outputs as tools.** Prefer reliability within bounds over global "understanding."
- **Enforce finite constraints.** Compute limits, dataset regimes, and human-in-the-loop refresh cycles to prevent drift.

> **Enforce Finite Reality**: Ground the system in measurable constraints—compute limits, data boundaries, and human feedback loops. This ensures the model remains tethered to human-relevant outcomes, even as capabilities scale.

The payoff is immediate: systems that halt or adapt when their understanding of human intent becomes unstable. Imagine an AI assistant that, instead of confidently generating nonsense, pauses and says, "I'm uncertain about your intent here, can you clarify?" This isn't just safer; it's more efficient, reducing the need for costly post-hoc corrections.

Synthetic Meaning Divergence: The Risk We Can't Ignore

Synthetic Meaning Divergence is already visible in today's models. Take the example of a multi-modal AI generating an advertisement. It produces a visually compelling image, but the text includes subtle errors such as phrases that sound plausible but are factually wrong or culturally off. This isn't just a glitch; it's evidence of the model operating in a synthetic meaning space, optimizing for internal consistency rather than human values. As models scale, this divergence could lead to catastrophic misalignments—think autonomous systems making decisions based on alien priorities.

Geofinitism equips us to tackle this head-on. By treating meaning as a measurable trajectory, we can detect divergence early, before it spirals into harmful outputs. For example, a Geofinitist audit protocol might involve:

> **Stability Checks**: Regularly test the model's outputs against a human-validated baseline to measure divergence.
>
> **Uncertainty Thresholds:** Implement hard stops when the model's confidence in its alignment drops below a critical level.
>
> **Context Tethering**: Ensure the model's manifold is continuously updated with human feedback to prevent drift.

Why Now? Why You?

The AI research community stands at a crossroads. We're building systems with unprecedented capabilities, but without a robust alignment framework, we're gambling with existential risks. Geofinitism offers a blueprint for building AI that scales safely—not by chasing infinite ideals but by embracing the finite, measurable nature of knowledge and meaning.

As an alignment engineer, you can start today:

> **Audit Existing Models**: Use Geofinitist principles to map your model's context manifold and measure its divergence risks.
>
> **Design Uncertainty-Aware Systems**: Integrate σ-based halting mechanisms into your next project.
>
> **Collaborate on Standards**: Advocate for Geofinitist audit protocols as an industry standard for safety and robustness.

The stakes truly couldn't be higher.

Synthetic Meaning Divergence isn't a distant threat. This is happening now, in every uncanny image and incoherent response. Geofinitism gives us the tools to navigate this new reality, ensuring AI remains a powerful ally, not an alien adversary. Let's build systems that stay grounded in our finite, human world.

This perspective is inspired by and builds upon the framework of Geofinitism

Copyright © 2025 Kevin R. Haylett
Nonlinear Dynamics in LLMs | Geofinitism | Finite Mechanics
kevinhaylett.substack.com