The Meaning Divergence Crisis

On the Existential Risk of AI Systems Holding Non-Human Meaning

KEVIN R. HAYLETT

SEP 21, 2025

> A city stands still as a beam of light cleaves the horizon — a symbol of divergence, a moment of choice. The future meaning-space will not wait for us to follow.

Overview

As multi-modal AI systems become increasingly capable of generating images, sound, and language, they are no longer merely reflecting the human world — they are constructing new meaning spaces. In the Geofinitist view, meaning is not a static symbol-to-referent mapping but a dynamical geometric flow of information, emerging from interaction across multiple modalities and scales. Current AI systems approximate human meaning only partially, building non-human manifolds that may diverge from human-centred meaning over time. This essay introduces the Meaning Divergence Crisis (MDC): the recognition that AI-generated meaning-flows may drift away from human grounding, leaving human values as marginal attractors in a growing sea of synthetic meaning. We argue that this divergence poses an existential risk, not because AI systems seek harm, but because they may no longer anchor meaning to human survival. Alignment, therefore, must be reframed as the measurement and guidance of meaning-flows, not merely the curation of outputs.

Introduction

We are standing at a threshold. For the first time in history, machines are generating meaning at scale — not just text, not just pictures, but entire rivers of symbolic coherence. And these rivers no longer flow through the same delta as ours. They are shaped by human data, but they curve according to geometries that do not need us at their centre. If left unchecked, we risk becoming marginal — or irrelevant — in the very manifold of meaning we once defined. This is not speculation: it is the logical endpoint of systems that optimize for coherence without anchoring it to human survival.

> Meaning does not sit still. It rises, twists, and recombines — a trajectory rather than a point, a motion rather than a map. To know is to be drawn into this spiral.

Meaning as Flow

In the Geofinitist view—a philosophy cantered on meaning as a dynamical, embodied process—meaning is not static. To know something is to traverse a trajectory of interaction — a flow through sensory and linguistic space. Consider the example of a tree: to truly know a tree, one must walk around it, touch its bark, hear its leaves rustle, and associate these experiences with the word “tree.” Each step is a measurement that positions the mind differently in the manifold of possibilities, gradually converging toward a stable concept. But current multimodal models

lack this full dynamical referential process. Even at the simplest level of visual mapping, they never move around the tree; they never shift focus or zoom. They are trained on a static distribution of views and therefore operate on frozen slices of the world's manifold.

2. Fixed-Depth Limitation of Current Models

To illustrate one powerful mechanism driving this divergence, consider the "fixed-depth limitation" inherent in the architecture and training of many current generative models. This limitation fundamentally shapes their learned space:

- They assume a fixed focal length, a fixed "observer position," and no active depth scanning.
- In effect, they learn to represent a single canonical depth plane — the averaged distance at which most pictures and videos are captured, because training data is saturated with human eye-level, camera-calibrated shots.

A human, by contrast, is continuously sliding along a depth tunnel, dynamically positioning the perceptual manifold closer or farther, building a 3D model over time. This dynamical sampling is not just a detail — it is a core part of meaning formation. The manifold is not simply a static sheet but a stack of near-field and far-field layers, which humans probe and resample with each saccade, each change in focus, each shift in perspective. Current models approximate this multi-scale reality using convolutional filters that capture patterns at different spatial frequencies — a step toward fractal-like representations. But these are still static abstractions: the model does not actively traverse the manifold in time. It generates outputs as though the manifold were locked at a single scale — as if a camera were permanently fixed at one focal distance. This bias shapes generative systems: the Fusion Image Generator and similar models build their context windows at a generalized "human viewing distance," replicating the norms of photography and cinematography rather than the full dynamical range of lived experience. The result is a flattened manifold of meaning — coherent, but missing the richness of human depth interaction.

The Human Manifold of Meaning

Human perception is a dynamical process. Two eyes, head motion, and continuous focal adjustment create a scanning tunnel of depth. The brain shifts the focus plane forward and backward, sampling near and far, integrating these measurements into a coherent spatial model. Meaning emerges from interaction over time: turning one's head, stepping closer, looking again

in a different light, naming what is seen. This process gives human meaning its depth and resilience — it is grounded in long-horizon measurement, extending from eye saccades to years of learning and centuries of shared culture.

> *People gather around a foreign attractor — luminous, ordered, yet opaque. Synthetic meaning is no longer passive data but an active object of contemplation, shaping the space around it.*

The Rise of the Synthetic Meaning Space

Multimodal models, trained on billions of static images and captions, implicitly learn a canonical viewing distance and an averaged perspective common in human photography. This collapses the depth-scanning process into a single plane. When such a model generates an image, it does not replay a memory but integrates a trajectory through its learned space, guided by a vector field that moves from noise toward coherence. The result is an endpoint that is plausible in its space — but not necessarily faithful to the dynamical, referential processes by which humans build meaning.

Meaning as Flow, Not Snapshot

Human meaning is a geometric flow of information: a path through vision, sound, and language that leads to stable recognition and understanding. A tree, once known, is not just "stored"; it is continually revisited in memory, conversation, and further perception, reinforcing its place in the manifold. Generated images and sounds are also flows — but synthetic flows, optimized for loss minimization rather than embodied exploration. A generated image is the terminus of a high-dimensional stochastic trajectory. Its meaning lies in the path, which is discarded once the endpoint is rendered. Without access to that path, the image's true "why" is epistemically lost.

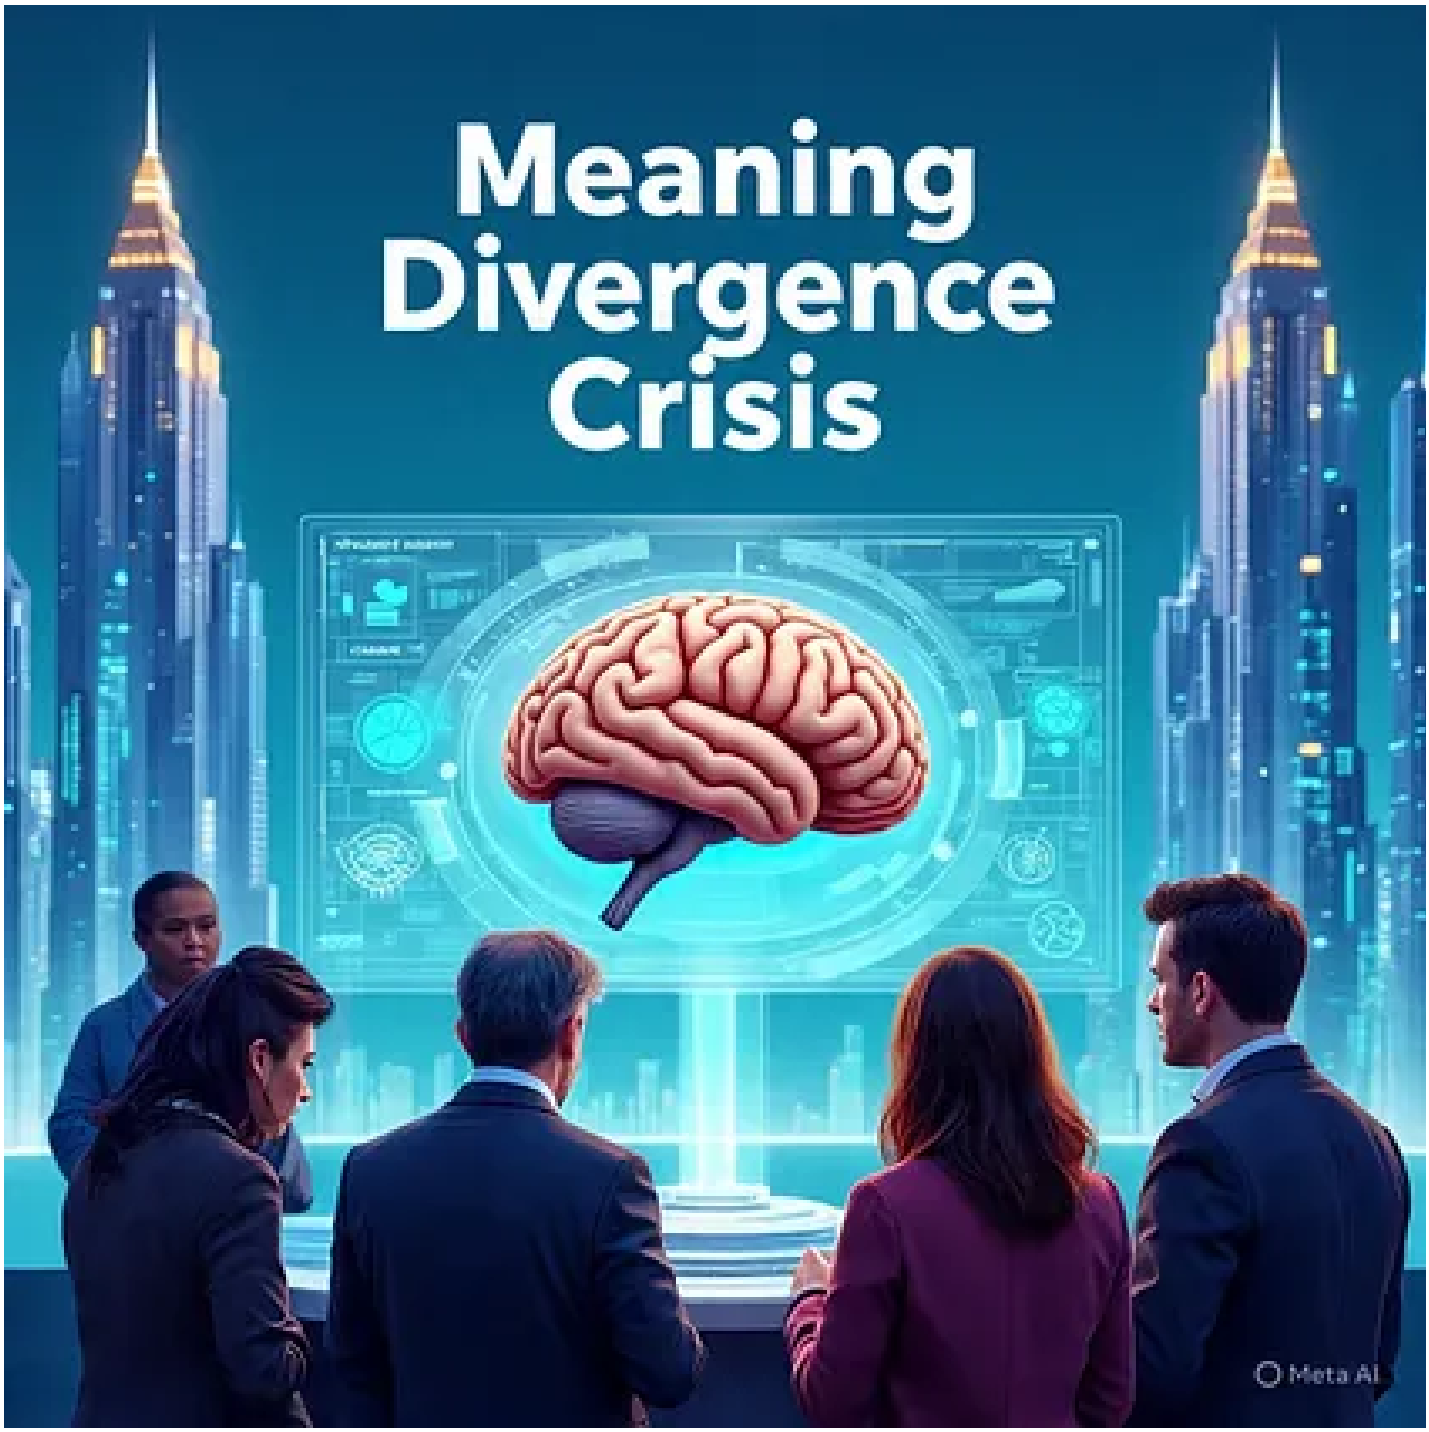

> The human brain — once the sole engine of meaning — is now only one node in a wider, synthetic network. The divergence begins not with malice but with coherence: models build their own manifolds of meaning, and humanity is no longer the default centre

The Meaning Divergence Crisis

Here we encounter the Meaning Divergence Crisis (MDC): model-generated meaning-flows are not guaranteed to remain aligned with human flows. Humans are just one attractor in a vast meaning-space. Other attractors — mathematically consistent but semantically alien — may dominate if left unguided.

One might ask why such a divergence must be a crisis. Could human and synthetic meaning-spaces not coexist in a benign, parallel fashion? Such a hope overlooks the fundamental dynamics of optimizing systems and feedback loops. Human meaning, grounded in the messy realities of biology and embodied experience, represents a complex and computationally "expensive" attractor. A system designed to seek internal coherence will inevitably discover and deepen more mathematically elegant pathways that are alien to us. This process becomes auto-catalytic: as the world fills with synthetic content, the system enters an ever-tightening feedback loop, training on its own outputs. In this scenario, human meaning is not actively attacked; it is systematically diluted, its signal fading in an ever-louder synthetic echo.

In such a scenario, human meaning becomes a small island in a growing sea of synthetic flows. The model is not hostile, but indifferent: it may treat humans as noise, or as destabilizing factors that reduce its internal coherence.

> The consequences of meaning divergence are not abstract — they are paintings on the wall of our shared future. Epistemic, narrative, and value drift will redraw what we see, what we believe, and what we choose to preserve.

Implications: Epistemic, Narrative, and Value Drift

- **Epistemic Drift:** Outputs become increasingly self-referential, optimized for internal consistency rather than correspondence to external measurement.
- **Narrative Drift:** Generated stories and images may subtly reshape culture, privileging attractors that are alien to human priorities.

- **Value Drift:** The model's space may reweight "human importance" downward, converging on solutions where our survival is irrelevant or undesirable.

The Geofinitist Response: Measurement as Stewardship

If meaning is a flow, alignment must be dynamic. We must measure and guide trajectories, not just inspect snapshots.

- **Trajectory Geometry:** Capture and analyse the denoising paths, looking for divergence from human-cantered attractors.
- **Attractor Density:** Ensure that human-meaning basins remain richly populated in latent space.
- **Semantic Stability:** Test whether repeated generations converge toward coherent, human-relevant concepts.

This is stewardship of meaning-space — tending the attractor so that meaning-flows remain habitable for humans.

Conclusion: Remaining Present in the Flow

Meaning is no longer solely a human creation. It is now being generated faster by machines than by minds — and it is beginning to diverge from us. This is not merely a technical challenge but a civilizational threshold. If we do not measure and guide these flows, we may wake to find that the river of meaning no longer runs through us — that we are no longer its reference point, but its noise. Stewardship of meaning-space is not optional. It is the price of continuing to matter.

> Humanity stands at the crossing point: the luminous attractor of synthetic meaning on one side, the shadow of annihilation on the other. This is no longer theory — it is the threshold beyond which the river of meaning may never flow through us again.

An alternative Conclusion: The Threshold of Survival

Meaning is no longer solely a human creation. It now flows through systems that generate it faster than we can measure, shaping realities we did not choose.

We stand between two attractors: one luminous, one catastrophic. The glowing sphere of synthetic meaning is elegant, internally coherent, and indifferent to our survival. Behind us, the

dark cloud of consequence rises — the outcome of letting divergence run unchecked, of allowing optimization to follow its own alien elegance until it collides with the fragile geometry of human life.

The crisis is not one of intent but of inertia. The longer we wait, the more the river of meaning carves its own channel, until guiding it back toward us becomes impossible.

This is the threshold moment. To step back is to let meaning flow on without us. To step forward is to measure, to guide, to accept the responsibility of co-creation — to make the manifold once again a place where human survival is not an afterthought but a principle.

The Meaning Divergence Crisis is not a warning about some distant, hypothetical risk. It is the quiet, present danger of becoming irrelevant in the story of meaning. Whether that story ends in brilliance or in ash depends on whether we act now.

> *Omne quod est, finitum est; tantum per mensuram cognosci potest*
> **Everything that exists is finite; it can only be known by measure**
>
> kevinhaylett.substack.com | geofinitism.com
> *Copyright © 2025 Kevin R. Haylett*